

Deep learning models for unbiased sequence-based PPI prediction plateau at an accuracy of 0.65

Timo Reim^{1,2} , Anne Hartebrodt² , David B. Blumenthal² , Judith Bennett^{1,*,†} , Markus List^{1,3,*} 

¹Data Science in Systems Biology, TUM School of Life Sciences, Technical University of Munich, Freising, 85354, Germany

²Department Artificial Intelligence in Biomedical Engineering, Friedrich-Alexander-Universität Erlangen-Nürnberg, Erlangen, 91052, Germany

³Munich Data Science Institute (MDSI), Technical University of Munich, Garching bei München, 85748, Germany

*Corresponding authors. Judith Bennett. Data Science in Systems Biology, Technical University of Munich, Maximus-von-Imhof Forum 3, Freising, 85354, Germany. E-mail: judith.bennett@tum.de; Markus List. Data Science in Systems Biology, Technical University of Munich, Maximus-von-Imhof Forum 3, Freising, 85354, Germany. E-mail: markus.list@tum.de.

[†]= equal contribution.

Abstract

Motivation: As most proteins interact with other proteins to perform their respective functions, methods to computationally predict these interactions have been developed. However, flawed evaluation schemes and data leakage in test sets have obscured the fact that sequence-based protein–protein interaction (PPI) prediction is still an open problem. Recently, methods achieving better-than-random performance on leakage-reduced PPI data have been proposed.

Results: Here, we show that the use of ESM-2 protein embeddings explains this performance gain irrespective of model architecture. We compared the performance of models with varying complexity, per-protein, and per-token embeddings, as well as the influence of self- or cross-attention, where all models plateaued at an accuracy of 0.65. Moreover, we show that the tested sequence-based models cannot implicitly learn a contact map as an intermediate layer. These results imply that other input types, such as structure, might be necessary for producing reliable PPI predictions.

Availability and implementation: All code for models and execution of the models is available at https://github.com/daisybion/PPI_prediction_study. Python version 3.8.18 and PyTorch version 2.1.1 were used for this study. The environment containing the versions of all other packages used can be found in the GitHub repository. The used data are available at <https://doi.org/10.6084/m9.figshare.21591618.v3>.

1 Introduction

Proteins perform a wide array of biological functions, but more than 80% depend on protein–protein interactions (PPIs) (Zhou *et al.* 2016). A comprehensive understanding of PPIs could significantly improve our knowledge of biological mechanisms, facilitate the identification of therapeutic targets, and aid in designing drugs for modulating specific protein interactions (Rao *et al.* 2014). Numerous experimental techniques, such as yeast two-hybrid, tandem affinity purification-mass spectrometry, X-ray crystallography, and NMR spectroscopy, have been developed to identify PPIs (Howell *et al.* 2006, Rao *et al.* 2014). However, they are cost-intensive and low-throughput compared to computational approaches (Rao *et al.* 2014). Since PPIs inherently occur in a three-dimensional context, including structural data would be ideal for this task. Unfortunately, obtaining 3D structures remains challenging despite the availability of in silico prediction tools like AlphaFold (Jumper *et al.* 2021). Hence, predicting PPIs based solely on amino acid sequences, which are more accessible, could provide an efficient alternative.

In a recent study, we highlighted issues in sequence-based PPI prediction methods, where improper data splitting introduced data leakage, inflating reported accuracies to levels exceeding 90% (Bennett *et al.* 2024). The leakage resulted from sequence similarity and protein node degree shortcuts between

the training and testing sets. Re-evaluating these models on a leakage-reduced gold-standard dataset (Bennett 2022) revealed significantly lower and often close to random performances, emphasizing the need for careful dataset preparation for unbiased method evaluation.

Since the identification of these issues, two notable models have been developed using our gold-standard dataset. Sledzieski *et al.* (2024) employed parameter-efficient fine-tuning of Evolutionary Scale Modeling-2 (ESM-2) embeddings. They achieved the best results with a baseline fully-connected neural network (FCNN) (accuracy around 0.63) from which they concluded that the lack of large datasets is currently the bottleneck in PPI prediction. Further, they found that training the FCNN on the smaller t33 embedding (650 million parameters) yielded better results than using the larger embeddings. Another model, TUNA, introduced by Ko *et al.* (2024a), applies transformer encoders followed by spectral normalized neural Gaussian processes for uncertainty-aware predictions. This architecture integrates spectral normalization with a Gaussian process to improve uncertainty estimation. They reached an accuracy of approximately 0.65. Additionally, three studies were published, enhancing the ESM-2 protein embeddings and evaluating their performance on our gold-standard dataset using only simple ML baselines. ProteinCLIP (Wu *et al.* 2024) combines ESM-2 and ProtT5

per-protein embeddings with natural language embeddings of protein function by OpenAI's text-embedding-3-large model. They achieve 0.61 to 0.65 accuracy. NaderiAlizadeh and Singh (2025) argue that dimension-wise average pooling to go from a per-residue to a per-protein representation is suboptimal as positions should not be weighted equally. They introduce a new pooling approach based on optimal transport. With both the ESM-2 average pooling approach and their new approach, they achieve about 0.65 accuracy. PoolPartI (Tartici *et al.* 2024) computes per-protein embeddings by combining the internal attention matrices of the ESM-2 model using PageRank and also reaches about 0.65 accuracy.

The fact that all of these approaches reach similar accuracies of around 0.65 with very different architectures and complexity levels raises a simple question: Are models for unbiased sequence-based PPI prediction attainable that yield accuracies that are substantially better than 0.65? Or is it likely that the field will plateau at this performance level, which would imply that focusing on models relying on other input types (e.g. structures) is necessary for reliably predicting PPIs?

In this work, we sought to answer this question by systematically assessing the effect of various techniques used in recently proposed PPI prediction models (overview in Fig. 1). Specifically, the state-of-the-art protein language model embeddings ESM-2 (Lin *et al.* 2023), as well as self and cross-attention modules, were assessed to gauge their individual and combined contribution to prediction performance. Previously published models such as Richoux (Richoux *et al.* 2019) and D-SCRIPT (Sledzieski *et al.* 2021) were modified using ESM-2 embeddings which led to notable improvements. The impact of model architecture and complexity was analyzed, comparing FCNNs, convolutional neural networks (CNNs), and transformer encoder-based models. Results indicate that embedding quality plays a more critical role in performance than model architecture, as no model achieved an accuracy higher than 0.65. Lastly, we investigated the possibility of implicitly predicting contact maps of interacting proteins as a penultimate layer, as this approach has been previously suggested (Sledzieski *et al.* 2021).

2 Methods

2.1 Data

Our previously introduced gold-standard dataset (Bennett 2022) was used for training, validation, and testing. This dataset is divided into three groups, Intra0, Intra1, and Intra2, with 59 260, 163 192, and 52 048 entries, respectively, each with an equal number of positive and negative interactions. The positive interactions stem from the HIPPIE v2.3 database (Alanis-Lobato *et al.* 2017). Intra1 is used for training, Intra0 for validation, and Intra2 for testing. The dataset is strongly leakage-reduced, since no protein is present in more than one split and sequence similarity between the splits is minimized (we used KaHIP (Sanders and Schulz 2013) to partition proteins based on sequence similarity and further reduced redundancy with CD-HIT at 40% pairwise sequence similarity). Since proteins have approximately the same node degree with respect to positive and negative edges, node degree bias is also reduced (Supplementary Methods, Figs S1–S3).

2.2 Protein embeddings

We generated protein embeddings from the amino acid sequences using ESM-2 (Lin *et al.* 2023), which then served as the input for the models. Embeddings from the three largest ESM-2 models, with dimensions d_{emb} of 5120, 2560, and 1280 were employed (referred to as t48, t36, and t33 with 15, 3, and 650 million parameters, each). The per-protein embeddings were obtained by dimension-wise average pooling of the per-token embeddings, resulting in a single vector of size d_{emb} for each protein. For models requiring per-token embeddings, proteins longer than 1000 amino acids were excluded due to computational limitations, leaving 93 719, 46 421, and 41 100 samples in the training, validation, and test sets, respectively.

2.3 Models

If not specified otherwise, the training set is always used for training, the validation set for hyperparameter tuning, and the test set for the final evaluation of the models. To counteract overfitting, a random subset containing 50% of the samples of the original training set is used to train the models in each epoch. All deep learning models use binary cross entropy (BCE) as the loss function and Adam for optimization. For models using unpadded per-token embeddings, each sample was processed separately, but backpropagation was performed for the entire batch simultaneously. We reimplemented the D-SCRIPT, Richoux, and TUnA models according to published method descriptions and modifications of the available code that were necessary to adapt the methods for training. Accordingly, even the vanilla reimplementations of D-SCRIPT, Richoux, and TUnA without algorithmic modifications may slightly differ from the model implementations underlying the original publications. Nonetheless, our implementation of Richoux and D-SCRIPT achieved approximately the same results on our gold-standard dataset (0.53 versus 0.52 and 0.5, respectively, Supplementary Table S2) as in our previous study, where we used their original implementations with the input representations from their respective publications (one-hot encoding for Richoux, Bepler and Berger (2019) embeddings for D-SCRIPT, more details in the Supplementary Methods). We also evaluated the performance of the other models when supplied with Bepler and Berger embeddings or one-hot encodings (Supplementary Table S2). We only slightly modified TUnA and reached the same accuracy as the original publication reached on our gold standard (Table 1).

2.3.1 Baseline models

We created two baseline models, one each for per-token and per-protein mean embeddings. The baseline models serve as a reference point for assessing the performance of the more complex models. They provide the minimal expectation that any complex model should surpass to be considered effective. For this reason, hyperparameter optimization was not performed on the baselines.

2.3.1.1 Random forest classifier baseline

A random forest classifier (RFC) provides a simple baseline for the one-dimensional per-protein embeddings. The RFC constructs 100 trees, each receiving a random subset of $\sqrt{2 \cdot f_{in}}$ features, with f_{in} being the total number of input features from one protein. The RFC was trained on the entire concatenated embeddings of both proteins as input. We also tested how much information can be extracted from a dimensionality-reduced ESM-2 embedding. For this, we used the first 200 or 20 components of a principal component




Figure 1. Overview of the analyses. (a) We computed ESM-2 embeddings of different sizes for the proteins of our data-leakage-reduced PPI dataset. The per-token embeddings have variable sizes depending on the protein length, while the per-protein embeddings have a fixed size by applying dimension-wise averaging. (b) We tested two models operating on the per-protein embeddings—a baseline random forest classifier and an adaption of the previously published Richoux model. Five models operated on the per-token embeddings: a 2d-baseline, the 2d-Selfattention and 2d-Crossattention models (which expanded the 2d-baseline through a Transformer encoder), and adaptions of the published models D-SCRIPT and TUNA. (c) Hyperparameter tuning gave us insight into the influence of each tunable parameter on the classification performance. (d) No model surpassed an accuracy of 0.65. The more advanced models had similar accuracies, suggesting that the information content of the ESM-2 embedding has more influence than the model architecture. Per-token models did not consistently outperform per-protein models. (e) We applied various modifications to test their influence: different embedding sizes, inserting a Transformer encoder into different positions, adding spectral normalization after the linear layers, self- versus cross-attention-attention, and removing the padding. (f) Finally, we compared the implicitly predicted distance maps of the 2d-baseline, 2d-Selfattention, 2d-Crossattention, and D-SCRIPT-ESM-2 to real distance maps computed from PDB structures. Created in BioRender.com.

analysis (PCA) of the embeddings. The RFC was trained on the concatenated vectors with either 400 or 40 features. The PCA was applied to the entire dataset of mean embeddings. This was done for each of the three embedding sizes. We refer to these models as “RFC-400” and “RFC-40”. We term the unreduced version “RFC-mean”.

2.3.1.2 2d-baseline

Creating a simple baseline model with the two-dimensional per-token embeddings as input is not trivial, as the variable input size prevents using simple machine learning methods

such as RFCs or support vector machines. As a result, the 2d-baseline (Supplementary Fig. S6) is a more complex model but still simpler than the previously published models. It receives the input embeddings of both proteins of dimensions $(len(p_1), d_{emb})$ and $(len(p_2), d_{emb})$ and reduces the embedding dimension with three linear layers with rectified linear units (ReLU) activation functions. The sizes of these layers are $d_{emb}/2$, $d_{emb}/4$ and 64. Next, the outer product along the second dimension of both tensors is computed, resulting in a tensor of shape $(len(p_1), len(p_2), 64)$. A convolution with one output channel is applied to this, followed by a pooling

Table 1. Comparison of all tested models.

Model	Accuracy	Precision	Recall	F1	AUPR	Loss	Training time	Best epoch	Training time per epoch
Baseline models									
RFC-40	0.577	0.629	0.374	0.496	0.584	—	474	—	—
2d-baseline	0.575	0.630	0.373	0.468	0.635	0.738	30845	22	1234
Attention models									
2d-Crossattention	0.641	0.660	0.589	0.623	0.678	0.643	29411	6	2101
2d-Selfattention	0.616	0.611	0.553	0.591	0.641	0.670	41326	14	1878
Adaptions of published models									
Richoux-ESM-2	0.633	0.627	0.654	0.640	0.655	0.653	2483	8	155
D-SCRIPT-ESM-2	0.628	0.638	0.602	0.619	0.636	0.650	29849	7	1990
TUnA	0.645	0.672	0.580	0.622	0.692	0.630	30013	7	2001

Accuracy, precision, recall, F1 score, AUPR (curves in [Supplementary Fig. S24](#)), and the BCE values summarized in the “Loss” column were computed on the test sets. The “Best epoch” column describes the training epoch with the best accuracy on the validation set, which is 8 epochs before training was terminated due to early stopping. Total training times and mean training times per epoch are reported in seconds. Best performances are indicated in bold.

operation. The maximum value of the resulting tensor is then transformed by the sigmoid function to form the final predicted label.

2.3.2 Attention models

To test the impact of self- or cross-attention, we created modified 2d-baselines ([Supplementary Fig. S7](#)). We added an encoder layer between the embedding reduction via linear layers and the outer product of both proteins and applied spectral normalization ([Supplementary Methods](#)) to all linear layers of the encoder. Both protein embeddings are fed separately through the same encoder. The self-attention model captures within-protein relationships; the cross-attention model focuses on connections between the two proteins. We also tested a version where the Transformer encoder was applied before the embedding reduction. We call these models “2d-Selfattention”, “2d-Crossattention”, “2d-Selfattention-encoder-pre-reduction”, and “2d-Crossattention-encoder-pre-reduction”. Models using encoder attention without spectral normalization have the suffix “no-spectral”.

2.3.3 Adaptions of published models

For the previously published models, hyperparameter optimization was performed using Weights & Biases (wandb) ([Biewald 2020](#)) with Bayesian optimization of the accuracy on the validation dataset. Optimizations were performed with different parameter combinations for each model, stopping the training process if the validation accuracy did not improve for eight epochs. All per-token models used exclusively the t33 embeddings of size 1280 due to computational limitations. The learning rate and the size of the first linear layer, which reduces the embedding size, were optimized for all models. If present, the dropout rate was also optimized. For models using convolutions and pooling, we optimized the kernel size and pooling type, while the number of heads for multi-head attention and the dimension of the FCNN in the encoder were optimized in the attention models. The exact hyperparameters and their associated range of values can be found in the code.

2.3.3.1 Adaptions of the Richoux model

The model proposed by [Richoux et al. \(2019\)](#) ([Supplementary Fig. S4b](#)) was one of the models with a near-random performance on the leakage-reduced gold-standard dataset introduced in our previous study, making it well-suited to show the impact of the ESM-2 embeddings and other changes. In contrast to the original Richoux model, our adaption receives the

per-protein embeddings from each protein as input. The embedding dimensions of each protein are reduced separately by two distinct linear layers. After concatenation, the two resulting vectors are processed through two additional linear layers. The size of the final layer is reduced to 1. From this single value, the predicted label is obtained from a sigmoid activation function. Every layer is followed by ReLU and batch normalization. The model was reimplemented with PyTorch instead of Tensorflow and we refer to it as “Richoux-ESM-2”. We also tested a version of the model in which the embeddings are first fed into a Transformer encoder, with and without adding spectral normalization (“Richoux-ESM-2-encoder(-no)-spectral”). Additionally, we tested applying spectral normalization after every linear layer (“Richoux-ESM-2-spectral”).

2.3.3.2 Adaptions of the D-SCRIPT model

The D-SCRIPT model by [Sledzieski et al. \(2021\)](#) was also tested in our previous study and including it here serves a similar comparative purpose as the Richoux-ESM-2 model. We reimplemented D-SCRIPT with PyTorch to ensure consistent testing (“D-SCRIPT-ESM-2”). It takes the variable-size per-token embeddings as input. The model ([Supplementary Fig. S4a](#)) reduces the embedding dimension of the per-token embeddings separately for each protein to size d , followed by a ReLU function and dropout layer. Next, it computes the absolute difference and element-wise product between each pair of elements in both transformed tensors. These are then concatenated along the last dimension, forming a new tensor of shape $(len(p_1), len(p_2), 2d)$. Using a convolution followed by batch normalization and ReLU, the embedding dimension is reduced to b . A matrix of shape $(len(p_1), len(p_2))$ is computed with another convolution, batch normalization, and ReLU. Sledzieski et al. treat this matrix as an implicitly predicted contact map of the two proteins. Finally, a custom interaction module developed by Sledzieski et al. is applied to generate the predicted label. In short, the interaction module applies a weighting to the matrix, performs max pooling, normalizes the result, and then uses a custom sigmoid function based on non-zero elements in the matrix. We also tested alternative D-SCRIPT adaptions, which included the insertion of a Transformer encoder before or after dimensionality reduction (“D-SCRIPT-ESM-2-encoder-pre-reduction”, “D-SCRIPT-ESM-2-encoder-post-reduction”). We also tested removing the spectral normalization after the encoder (“D-SCRIPT-ESM-2-encoder-pre-reduction-no-spectral”) and applying cross- instead of self-attention in the encoder

(“D-SCRIPT-ESM-2-encoder-cross”). The family of D-SCRIPT adaptions is called “D-SCRIPT-like”.

2.3.3.3 Adaptions of the TUnA model

TUnA was developed by Ko *et al.* (2024a) and already uses ESM-2 embeddings as protein representation in the original version. In this model (Supplementary Fig. S5), all input embeddings of proteins with length < 1000 are padded to size $(1000, d_{emb})$. The embedding dimension is first reduced by a linear layer to d_{att} before the Intra Encoder module is applied. This module utilizes self-attention on the single sequences to recognize important relationships of different positions in the same protein. Though both inputs are processed individually, the weights are shared. The updated tensors x_1 and x_2 are concatenated to a combined tensor x_{12} , which is then sent to the Inter Encoder module to capture important relationships between the two proteins with self-attention. To ensure permutation invariance, this is repeated for the combined tensor x_{21} . In both encoded tensors, the average over all non-padded sequence positions is computed for each embedding dimension, resulting in two vectors of length d_{att} . These are combined into a single vector of length d_{att} by taking the position-wise maximum. This interaction feature vector is then turned into a prediction by a random feature expansion of a Gaussian process (RFEGP) layer, followed by a sigmoid function. All linear layers, including those in the encoders, are spectral normalized. This, along with the RFEGP, serves to enhance the model’s uncertainty awareness while retaining accuracy. To test the influence of padding on the model performance, we implemented a version using unpadded per-token embeddings (“TUnA-unpadded”, for results, see Supplementary Material). We also tested removing the spectral normalization and using a cross-attention-attention mechanism in the encoders (“TUnA-no-spectral”, “TUnA-crossattention”). The family of TUnA adaptions is referred to as “TUnA-like”.

2.4 Distance map prediction

Further, we tested whether models with per-token embedding input can implicitly predict a distance map. This had first been suggested by Sledzieski *et al.* (2021) with their D-SCRIPT model. The underlying idea is to have an $n \times m$ tensor as a penultimate layer obtained through convolution operations of the higher-dimensional tensors. The model might, hence, implicitly learn important positions like binding sites. We implemented this architecture for the 2d-base-line, Self-, Crossattion, and D-SCRIPT-like models. We performed a case study investigating the differences between real, experimentally determined distance maps and the implicitly predicted distance maps. To acquire the real distance maps, we searched the Protein Data Bank (PDB) (Berman *et al.* 2000) for protein complexes containing only the two proteins of the interactions of the gold-standard dataset, as other proteins, cofactors, or ligands cannot be taken into consideration by the models. Complexes containing homomers were also ignored. We filtered out entries where structural information was available exclusively for small parts of the proteins, i.e., peptides stemming from the original protein. Still, the protein sequences from the PDB were mostly shorter than the sequences used as input for prediction. Additionally, many proteins had modifications at certain amino acids that were added in the experiment that generated the data. To account for both of these issues, local alignment was used to

match the sequences from the predictions to those of the experimental data. We used only those predicted distance maps where the associated prediction was above the threshold of 0.9 to ensure the models’ confidence in the prediction. As a numerical value for similarity between experimentally determined and predicted distance maps, we computed the Pearson correlation of the two matrices.

3 Results

3.1 Hyperparameter optimization did not yield better parameter combinations than defaults

All models, apart from the Richoux-ESM-2 model, ran with 40 different hyperparameter configurations. Due to its considerably shorter runtime, approximately 250 different combinations were tested for the Richoux-ESM-2 model. The initial goal of the optimization was to identify the configuration that produces the best results. However, no parameter combination led to superior performances than the configurations using the original hyperparameter values (for the D-SCRIPT-ESM-2, Richoux-ESM-2, and TUnA) or the parameters that we chose originally (for 2d-Selfattention and 2d-Crossattention). Nonetheless, the optimization produced valuable insights into the impact of specific parameters on the performance, as well as into the robustness of the models themselves. For Richoux-ESM-2, the hyperparameters had minimal impact on the performance (Supplementary Fig. S9). For the models using an attention mechanism (2d-Selfattention, 2d-Crossattention, TUnA), the learning rate was always reported to be the most important parameter with high negative correlation (Supplementary Figs S10–S12). Finally, the 2d-Selfattention model favored max pooling, while the 2d-Crossattention model produced slightly better results with average pooling.

3.2 ESM-2 embeddings boost test accuracies of all models to around 0.65

For the performance evaluation of all models on the test set, we used the best-performing hyperparameter combinations (Table 1). The exact configurations can be found in the code of the associated GitHub repository. Of note, this resulted in all per-token models using the ESM-2 t33 embeddings and Richoux-ESM-2 using t36 embeddings. The RFC baseline with the best performance was RFC-40 on the t36 embeddings (Fig. 2).

TUnA achieves the highest accuracy, precision, and lowest BCE loss, closely followed by the 2d-Crossattention model. Notably, all models, except Richoux-ESM-2, exhibit much higher precision than recall, leading to reduced F1 scores.




Figure 2. Validation performance with increasing ESM-2 embedding size. Most models perform best with the smaller t33 embedding.

This imbalance is more pronounced in models with attention mechanisms, suggesting a conservative approach that minimizes false positives but sacrifices recall by missing true positives. In contrast, Richoux-ESM-2 achieves a better balance between precision and recall, improving its detection of true positives without significantly increasing false positives. None of the models delivers a reliable PPI prediction, with accuracies not exceeding 0.65. All complex models fall within a narrow accuracy range of 0.61–0.65, regardless of their architectural differences. Using the ESM-2 embeddings, the adapted Richoux-ESM-2 and D-SCRIPT-ESM-2 now achieved accuracies of 0.633 and 0.628, demonstrating that the embeddings elevate models from near-random to near-state-of-the-art performance. Similarly, the RFC, previously achieving 0.53 accuracy, improves to 0.577, highlighting the embeddings' information-rich content. This consistency suggests the performance is primarily driven by the ESM-2 embeddings rather than by the models themselves. This observation is further underlined by our results in [Supplementary Table S2](#) (comparison ESM-2 embeddings versus Bepler and Berger embeddings and one-hot encoding).

3.3 Per-token embeddings do not yield better results than averaged per-protein embeddings

[Supplementary Table S1](#) gives an overview of all model modifications. The modifications were evaluated on the validation set, since we consider them part of the model design. While the embeddings substantially enhance performance, the choice of embedding type (mean or per-token) has minimal impact. The baselines show virtually no difference. The more advanced models have no visible trend. While the best-performing models, TUNA and 2d-Crossattention, are per-token models, Richoux-ESM-2, which uses the mean embeddings, outperforms D-SCRIPT-ESM-2 and 2d-Selfattention. Overall, the differences in performance are small. This finding is unexpected, as per-token embeddings should theoretically encode more interaction-specific information than mean embeddings. Regarding training, models with mean embeddings are significantly faster. Complex models employing per-token embeddings require around 2000 s per epoch, whereas the Richoux-ESM-2 model reduces this by a factor of 13 with comparable performance. RFC-40 achieves an additional order-of-magnitude speedup, though this excludes PCA runtime.

3.4 Models profit from smaller embeddings

[Sledzieski et al. \(2024\)](#) stated that using smaller input embeddings led to superior results in training their model. Although embedding size had little impact on Richoux-like models, others, such as TUNA, 2d-baseline as well as 2d-Self- and 2d-Crossattention showed increased performance on the t33 embeddings compared to the larger t36 and t48 embeddings ([Fig. 2](#)). The similar performance of the different embedding sizes on the training set could indicate a relationship between larger input embeddings and the susceptibility to overfitting, possibly due to the increase in model parameters.

3.5 Only the 2d-baseline profits from attention

The 2d-Selfattention and 2d-Crossattention models differ from the 2d-baseline only by a transformer encoder layer, which is inserted between the embedding reduction and the combined representation of both sequences ($\text{len}(p_1) \times \text{len}(p_2) \times 64$). This modification resulted in notable performance improvements. 2d-Crossattention outperforms 2d-Selfattention, likely due to

its ability to capture long-range protein interactions that complement the short-range relationships captured by convolution layers. Although the Transformer encoder nearly doubles runtime per epoch, the performance gains justify this trade-off.

To test the impact of the attention mechanism on published model architectures, we compared the regular D-SCRIPT-ESM-2 and Richoux-ESM-2 models to the versions utilizing an encoder (D-SCRIPT-ESM-2-encoder-post-reduction, Richoux-ESM-2-encoder-spectral). For Richoux-ESM-2-encoder-spectral, the attention mechanism led to a performance drop ([Fig. 3](#)). This drop is likely due to an incompatibility of mean embeddings and the attention mechanism. In self-attention, the goal is to capture relationships in the input by modifying the embedding space of the related positions accordingly. By taking the average feature values of all sequence positions, the related positions enhanced by the self-attention will lose any meaning. Furthermore, the average feature values are altered by the attention mechanism, possibly making them partially unreadable or uninterpretable.

For D-SCRIPT-ESM-2-encoder-post-reduction, the addition of the encoder had no notable influence on the model performance. This was surprising, given that the 2d-Selfattention and 2d-Crossattention models have a similar architecture to D-SCRIPT-ESM-2. One reason could be the performance barrier of 0.65 accuracy, which seems to be the information content of the ESM-2 embeddings. A model that can already harness the information contained in the embeddings and manages to approach the barrier by itself may not be able to benefit from the attention mechanism. A simple model that is not yet able to capture this information fully might be enhanced to do so through the attention mechanism.

As the 2d-Crossattention model outperforms the 2d-Selfattention model, we explored this aspect further by examining the influence of cross-attention-attention versus self-attention on the D-SCRIPT-ESM-2-encoder-post-reduction and TUNA models ([Fig. 3](#)). The D-SCRIPT model shows similar performance regardless of the type of attention mechanism, while TUNA-crossattention is slightly better. The results of TUNA-crossattention also question the meaningfulness of the TUNA architecture, as the idea of having an intra-protein self-attention encoder that captures relationships inside each of the proteins and an inter-protein self-attention encoder that captures relationships between the two proteins is similar to adding cross-attention-attention.

3.6 Attention-based models profit from spectral normalization

The TUNA publication argues that adding spectral normalization after linear layers increases model performance. Indeed, removing all spectral normalization layers in the TUNA model




Figure 3. Validation performances with and without the addition of a self-/cross-attention-attention encoder to the respective models.

resulted in random predictions (Fig. 4). Motivated by this behavior, we investigated the influence of adding spectral normalization to the other models. On Richoux-ESM-2, spectral normalization on the linear layers had no clear effect. On all attention-based models (2d-Selfattention, 2d-Crossattention, D-SCRIPT-ESM-2-encoder-post-reduction, and Richoux-ESM-2-encoder), the absence of spectral normalization resulted in random predictions (Fig. 4).

While investigating possible causes of this, we discovered that models using attention without spectral normalization stagnated during training. Irrespective of the input, the model went back and forth between either predicting all interactions as positive, with values slightly above 0.5, or negative, with values slightly below 0.5. We can only speculate about the reasons. It could be that, without spectral normalization, the gradients in the attention mechanism exponentially grow or decay, leading to nonsensical output. Interestingly, TUNA-like (Supplementary Fig. S13) even learns regularly for the first few epochs before also dropping to random performance. Accordingly, all comparisons of models using attention were done using spectral normalization.

3.7 Attention should be applied after reducing embedding size

We explored the influence of reducing the embedding size via linear layers prior to the attention mechanism or encoder. In all tested models, reducing embedding size leads to superior performance on the validation dataset (Fig. 5). D-SCRIPT-ESM-2-encoder-pre-reduction had near-random performance scores. We hypothesize that not reducing the embedding dimension may have a similar effect as using a larger embedding, which, as stated before, may be the cause for increased overfitting, explaining the performance of the 2d-Selfattention model. The diminished performance of both 2d-Crossattention-encoder-pre-reduction and D-SCRIPT-ESM-2-encoder-pre-reduction on the training and validation sets may also hint at the negative impact of using many input parameters for the attention mechanism, as it could potentially obscure critical positions among the noise.




Figure 4. Validation performance with and without removing spectral normalization after the linear layers for the models including an encoder and Richoux-ESM-2.




Figure 5. Validation performance for inserting the encoder before or after dimensionality reduction via linear layers.

3.8 Distance maps cannot be predicted implicitly

Of the tested models, D-SCRIPT-ESM-2, the 2d-baseline, 2d-Selfattention, and 2d-Crossattention internally form a matrix of size $\text{len}(p_1) \times \text{len}(p_2)$. As D-SCRIPT (Sledzieski et al. 2021) treats this as implicit contact map prediction, we investigated the similarities of this matrix to the real, experimentally determined contact maps. For this case study, we used all protein pairs of the dataset, regardless of the split they appeared in, as the models cannot overfit the distance maps during training because they received no structural information. Matching the confident predictions (> 0.9) with the PDB entries resulted in a total of 84 interactions across all models. Of those, 54 were filtered out due to them containing other cofactors, ligands, or homomers, and another 19 were removed because the protein sequences in the experimental data were too short. This left 3 interactions for D-SCRIPT-ESM-2, 1 for 2d-Crossattention, 7 for 2d-Selfattention, and none for the 2d-baseline (Supplementary Table S3).

In the predicted distance maps from D-SCRIPT-ESM-2, certain structures are formed, which is especially visible in Fig. 6. In Supplementary Fig. S14, some vertical lines from the real distance map can be found in the prediction, indicating that the model identified some of the important positions of the Q8NAV1 protein. In the 1B34 complex (Supplementary Fig. S15), there are several clusters of high values in the bottom left and clusters of low values in the top right corner. While these clusters do not align with any clear features of the real distance map, they do exhibit a slight negative correlation.

Like D-SCRIPT-ESM-2, the 2d-Selfattention model also confidently predicted the interaction of complex 1B34 (Supplementary Fig. S16). In contrast to D-SCRIPT-ESM-2, no clear, distinct clusters or features are formed. Rather, some horizontal lines can be identified in both distance maps, hinting at the successful identification of some important positions of P63214. While both models achieve the same, although inverse, correlation, the 2d-Selfattention prediction resembles a seemingly structureless grid rather than an actual distance map. The same grid can be observed in the predictions for 2V8S for both 2d-Self- and 2d-Crossattention. Interestingly, 2d-Selfattention more clearly detects the vertical lines (Supplementary Fig. S16), unlike the 2d-Crossattention model, which more distinctly recognizes the horizontal features (Supplementary Fig. S18). Similar results can be observed for the remaining maps predicted by 2d-Selfattention (Supplementary Figs S19–S23).

Ultimately, the observed low correlations and the notable disparity between the real and predicted maps show that no model is capable of indirectly predicting distance maps or even detecting more complex structural features. This was to be expected, considering that the models receive only sequential data and that the task of predicting 3D structure from sequence was only partially solved by AlphaFold (Jumper et al. 2021), which is considerably more complex. However, the models were able to identify some singular, important positions for the interaction.

4 Discussion

In this study, we comprehensively investigated the efficacy of PPI prediction of various deep learning models, including Richoux-ESM-2, D-SCRIPT-like, and TUNA-like models, as well as 2d-Selfattention, and 2d-Crossattention models. Hyperparameter optimization across 40 configurations had




Figure 6. Comparison of the real and predicted distance map of the D-SCRIPT-ESM-2 model for the 2A24 complex. Low values (i.e., contacts in the distance map) are represented by bright areas, while high values correspond to darker regions. The values of the implicitly predicted distance map correspond to the output of the final convolution layer of the D-SCRIPT-ESM-2 model followed by batch normalization and ReLU activation.

little impact, and all models struggled to surpass an accuracy of 0.65, implying that the performance may be dictated more by the embeddings than the architecture of the individual models. This is further supported by the literature, where similar results are reported. In terms of speed, models utilizing averaged per-protein embeddings proved faster than their per-token counterparts, while producing similar results. We also tested several modifications of the models. For the relatively simple 2d-Crossattention and 2d-Selfattention models, adding a Transformer encoder with an attention mechanism led to a notable increase in performance compared to the 2d-baseline. In contrast, performances did not improve for more complex models, which reached test accuracies close to 0.65 already without Transformer encoders. This further supports our hypothesis that this value constitutes a barrier that is unlikely to be surpassed by purely sequence-based models relying on ESM-2 embeddings.

In comparing real distance maps with those predicted by the models, we tested the hypothesis of implicitly predicting contact maps as suggested in the D-SCRIPT paper (Sledzieski *et al.* 2021). While the D-SCRIPT-ESM-2 model produced some recognizable structures in these maps, they showed little to no relation to the experimentally determined distance maps. Similar results were observed for the 2d-Selfattention and 2d-Crossattention models. These findings highlight the broader challenge of accurately predicting structural data in pairwise PPIs, given the limited availability of such data.

Despite these highly relevant findings, our study is subject to at least four limitations. First, we have only investigated ESM-2, the most commonly used embedding in the latest publications, and the less complex Bepler & Berger and one-hot embedding. Other embeddings like the newer, sequence-based ESM-C model (ESM Team 2024) or multimodal embeddings might push the observed upper bound. Second, many compromises were made due to computational limitations, such as restricting the input in embedding size and sequence length, as well as an inexhaustive optimization of model hyperparameters. For the same reason, we did not do robustness tests.

Further, the availability of PPI data for training, validation, and testing also poses an issue. While our gold-standard dataset is an improvement compared to previous PPI datasets, many top-performing ML models, such as AlphaFold (Jumper *et al.* 2021), are trained on substantially larger datasets.

When screening PDB for data on protein complexes for our comparison between experimentally determined and implicitly predicted distance maps, we identified only very few suitable complexes that contain only two proteins. The PDB reports 2836 pairwise protein interactions without any cofactors, ligands, or other proteins, while it reports 46 849 entries with at least one additional subunit. Furthermore, some interactions described as involving only two proteins may actually consist of homomers. This points to a general limitation of existing sequence-based PPI prediction models: While existing models are designed to predict PPIs consisting of only two proteins, most complexes formed in permanent interactions are comprised of many more proteins, other cofactors or ligands, which cannot be captured in pairwise interactions but may be essential for the interaction.

Another key challenge lies in the choice of protein embeddings, as both mean and per-token approaches present significant limitations for sequence-based PPI prediction. Mean embeddings lack positional information, making them unsuitable for capturing structural details such as interacting amino acids, especially when all positions are weighted equally during the average pooling. In contrast, per-token embeddings retain positional data but introduce variability in sequence length, requiring computationally expensive padding or alternative handling. Padding, while standard, can inflate computational costs and lead to uneven training due to sequence length distributions. Moreover, models trained on padded sequences may fail to process proteins longer than the training limit. Unpadded embeddings avoid these issues but restrict sequence length reductions to basic operations, limiting flexibility. Per-token embeddings may perform well for position-specific tasks (e.g. predicting the effects of single nucleotide polymorphisms), but their application to PPI is

constrained by the need for experimental identification of contact positions in unseen proteins.

Given these limitations, we suggest the following directions for future work: First, enhancing the embeddings with non-sequential data might boost the performance. Fan *et al.* (2025) distinguish between multiple sequence-based embeddings, structure- and function-enhanced embeddings, and multimodal approaches, including text and molecule information. The Bepler & Berger embedding, which includes structural information, does not perform much worse than ESM-2 despite a less complex architecture and seeing fewer proteins during pretraining (Supplementary Table S2). Ko *et al.* (2024b), who fuse several protein-text embeddings, reach an accuracy of 0.68 on our dataset. ESM-3 (Hayes *et al.* 2025) tackles the challenge of the currently scarce suitable structural data by incorporating predicted structures (as well as functional annotations). Incorporating genomic information like gLM2 (Cornman *et al.* 2024) may also be interesting. Predicting contact regions or specific interaction sites could further refine inputs, reducing noise and enhancing the utility of per-token embeddings.

Author contributions

T. Reim (Formal analysis [lead], Investigation [lead], Methodology [lead], Software [lead], Validation [lead], Writing—original draft [lead]), A. Hartebrodt (Project administration [equal], Writing—review & editing [equal]), D. B. Blumenthal (Funding acquisition [equal], Project administration [equal], Writing—review & editing [equal]), J. Bennett (Conceptualization [equal], Data curation [lead], Project administration [equal], Supervision [equal], Visualization [lead], Writing—review & editing [equal]), M. List (Conceptualization [equal], Funding acquisition [equal], Project administration [equal], Resources [lead], Supervision [equal], Writing—review & editing [equal])

Supplementary data

Supplementary data are available at *Bioinformatics* online.

Conflict of interest: D.B.B. consults for BioVariance GmbH. M.L. consults for mbiomics GmbH. All other authors declare no competing interest.

Funding

T.R., A.H., and D.B.B. were supported by the German Federal Ministry of Education and Research (BMBF) within the framework of the CompLS funding concept [031L0309A (NetMap)]. M.L. and D.B.B. were supported by the Klaus Tschira Stiftung [00.003.2024]. J.B. and M.L. were supported by the German Federal Ministry of Education and Research (BMBF) within the framework of the CompLS funding concept [031L0305A (DROP2AI)]. Funded by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation) [422216132].

Data availability

All code for the models and execution of the models is available at https://github.com/daisybio/PPI_prediction_study. The gold-standard dataset is available in figshare, DOI: 10.6084/m9.figshare.21591618.v3.

References

- Alanis-Lobato G, Andrade-Navarro MA, Schaefer MH. HIPPIE v2.0: enhancing meaningfulness and reliability of protein-protein interaction networks. *Nucleic Acids Res* 2017;45:D408–D414.
- Bepler T, Berger B. Learning protein sequence embeddings using information from structure. arXiv [cs.LG], <https://doi.org/10.48550/arXiv.1902.08661>, 2019, preprint: not peer reviewed.
- Berman HM, Westbrook J, Feng Z *et al.* The protein data bank. *Nucleic Acids Res* 2000;28:235–42.
- Bennett J. PPI prediction from sequence, gold standard dataset. Figshare, <https://doi.org/10.6084/m9.figshare.21591618.v3>, 2022.
- Bennett J, Blumenthal DB, List M. Cracking the black box of deep sequence-based protein–protein interaction prediction. *Brief Bioinform* 2024;25:bbae076.
- Biewald L. Experiment tracking with weights and biases. 2020. <https://www.wandb.com/>. Software available from wandb.com.
- Cornman A, West-Roberts J, Camargo AP *et al.* The OMG dataset: an Open MetaGenomic corpus for mixed-modality genomic language modelling. *bioRxiv*, <https://doi.org/10.1101/2024.08.14.607850>, 2024, preprint: not peer reviewed.
- ESM Team. ESM Cambrian: revealing the mysteries of proteins with unsupervised learning. EvolutionaryScale Website, 2024. <https://evolutionaryscale.ai/blog/esm-cambrian> (26 April 2025, date last accessed).
- Fan W, Zhou Y, Wang S *et al.* Computational protein science in the era of large language models (LLMs), arXiv [cs.CE]. <https://doi.org/10.48550/arXiv.2501.10282>, 2025, preprint: not peer reviewed.
- Hayes T, Rao R, Akin H *et al.* Simulating 500 million years of evolution with a language model. *Science* 2025;387:eads0018–858.
- Howell JM, Winstone TL, Coorssen JR *et al.* An evaluation of in vitro protein–protein interaction techniques: assessing contaminating background proteins. *Proteomics* 2006;6:2050–69.
- Jumper J, Evans R, Pritzel A *et al.* Highly accurate protein structure prediction with AlphaFold. *Nature* 2021;596:583–9.
- Ko YS, Parkinson J, Liu C *et al.* TUmA: an uncertainty-aware transformer model for sequence-based protein–protein interaction prediction. *Brief. Bioinform* 2024a;25:bbae359.
- Ko YS, Parkinson J, Wang W. Benchmarking text-integrated protein language model embeddings and embedding fusion on diverse downstream tasks. *bioRxiv* 2024b. <https://doi.org/10.1101/2024.08.24.609531>
- Lin Z, Akin H, Rao R, Hie B *et al.* Evolutionary-scale prediction of atomic level protein structure with a language model. *Science* 2023;379:6637.
- NaderiAlizadeh N, Singh R. Aggregating residue-level protein language model embeddings with optimal transport. *Bioinformatics Advances* 2025;5:1.
- Rao VS, Srinivas K, Sujini G *et al.* Protein–protein interaction detection: methods and analysis. *Int J Proteomics* 2014;2014:147648.
- Richoux F, Servantie C, Borès C *et al.* Comparing two deep learning sequence-based models for protein–protein interaction prediction. arXiv [cs.LG], <https://doi.org/10.48550/arXiv.1901.06268>, 2019, preprint: not peer reviewed.
- Sanders P, Schulz C. Think locally, act globally: highly balanced graph partitioning. In: SEA 2013, LNCS, Vol. 7933, 2013, 164–175.
- Sledzieski S, Singh R, Cowen L *et al.* D-SCRIPT translates genome to phenotype with sequence-based, structure-aware, genome-scale predictions of protein–protein interactions. *Cell Syst* 2021;12:969–82.e6.
- Sledzieski S, Kshirsagar M, Baek M *et al.* Democratizing protein language models with parameter-efficient fine-tuning. *Proc Natl Acad Sci U S A* 2024;121:e2405840121.
- Tartici A, Nayar G, Altman RB. Pool PaRTI: a PageRank-based pooling method for robust protein sequence representation in deep learning. *bioRxiv*, <https://doi.org/10.1101/2024.10.04.616701>, 2024, preprint: not peer reviewed.
- Wu KE, Chang H, Zou J. ProteinCLIP: enhancing protein language models with natural language. *bioRxiv*, doi: [10.1101/2024.05.14.594226](https://doi.org/10.1101/2024.05.14.594226), 2024, preprint: not peer reviewed..
- Zhou M, Li Q, Wang R. Current experimental methods for characterizing protein–protein interactions. *ChemMedChem* 2016;11:738–56.

© The Author(s) 2025. Published by Oxford University Press.

This is an Open Access article distributed under the terms of the Creative Commons Attribution License (<https://creativecommons.org/licenses/by/4.0/>), which permits unrestricted reuse, distribution, and reproduction in any medium, provided the original work is properly cited.

Bioinformatics, 2025, 41, 590–598

<https://doi.org/10.1093/bioinformatics/btaf192>

ISMB/ECCB 2025 Supplement